

Module 2 : Core Services



Amazon Virtual Private Cloud (VPC)

Introduction

AWSSOME DAY
ONLINE CONFERENCE

Private, virtual network in the AWS Cloud
Similar constructs as on-premises network
Complete control of network configuration

Features

Characteristics

- Allows you to provision virtual networks

Logically isolated

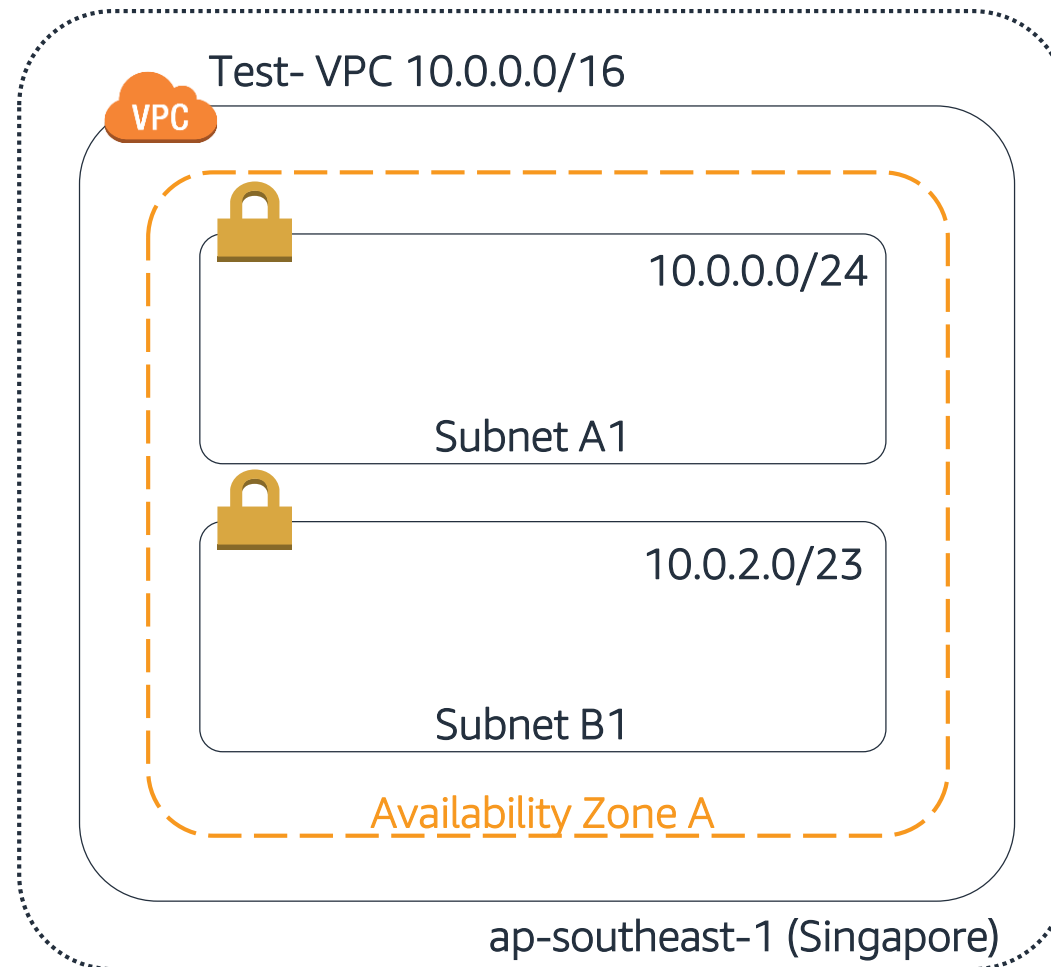
Configurable key features

- IP ranges
- Routing
- Network gateways
- Security settings

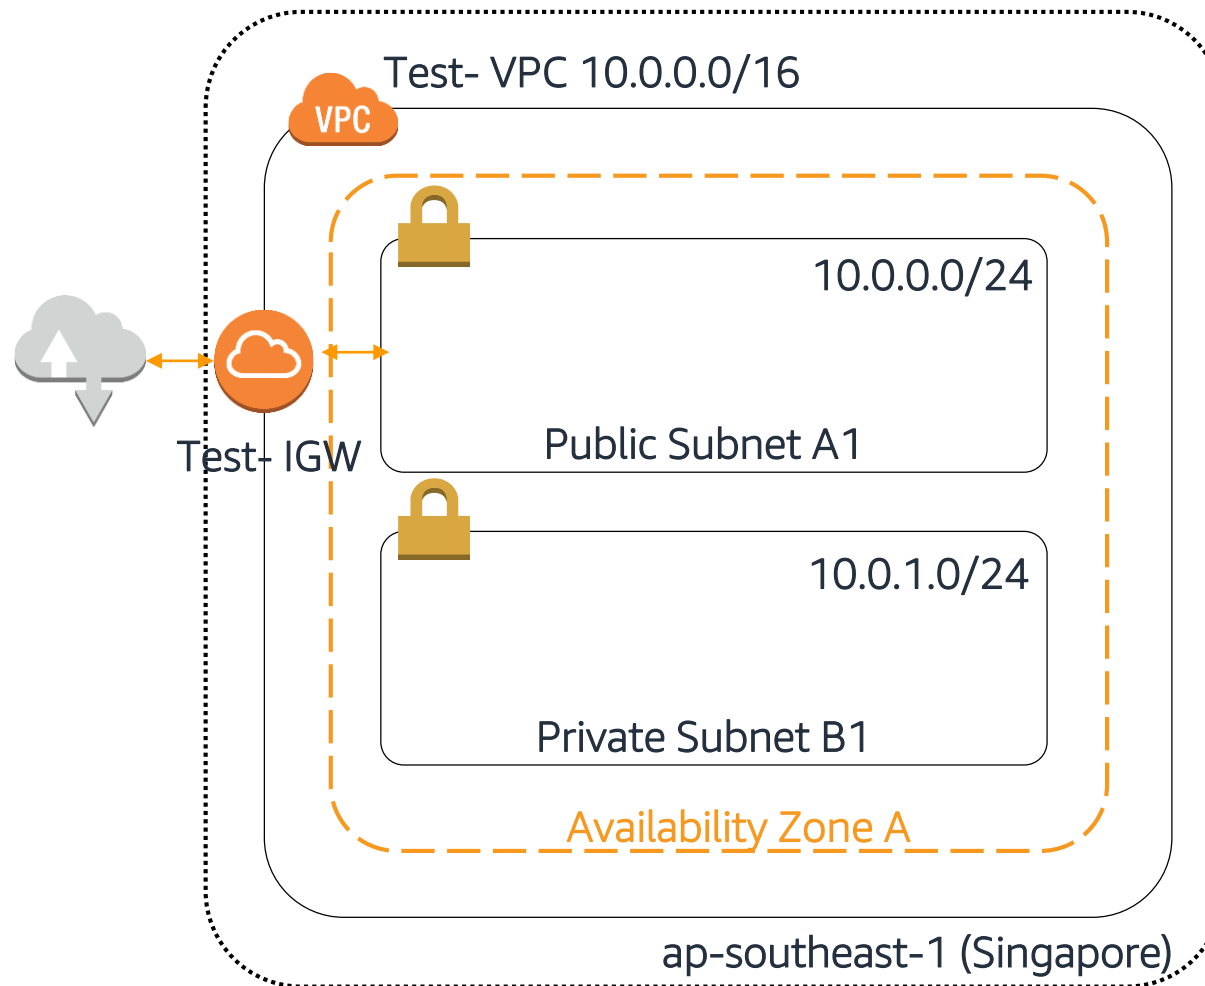
Route Tables

- Control traffic going out of the subnets

Example



Example



VPC concepts:

- Region Scope
- An internet gateway
- Public subnet
- Private subnet

Learn More

- Route tables and isolation methods
- Other Amazon VPC features (e.g., VPC endpoints and VPC peering connections)

AWS Security Groups

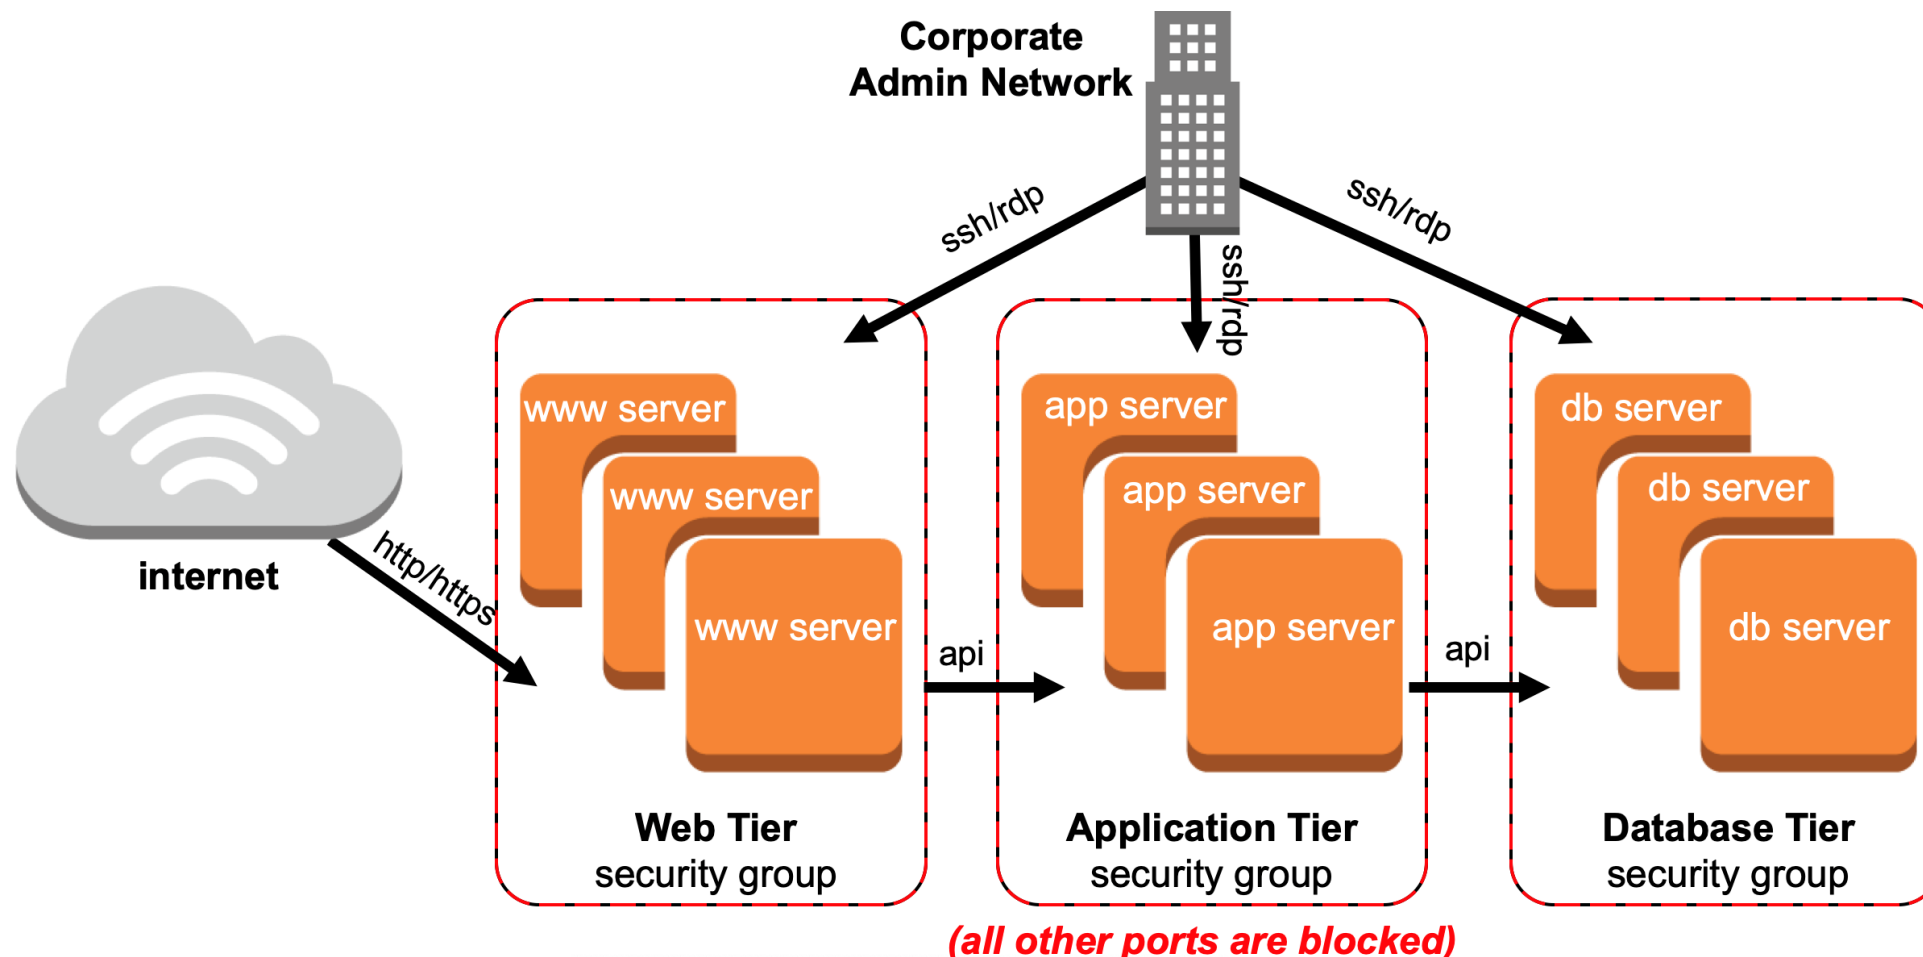
Is the highest priorities Security groups

- Act as built-in firewalls
- Control accessibility to instances



AWS Security

AWSSOME DAY
ONLINE CONFERENCE



Compute Services

Compute Services

AWSSOME DAY
ONLINE CONFERENCE

AWS

- Flexible
- Cost effective

Amazon EC2

- Flexible configuration and control

AWS Lambda

- Pay only for what you use
- No administration

Compute Services

AWSSOME DAY
ONLINE CONFERENCE

Amazon Lightsail

- Launch virtual private server
- Manage simple web and application servers

Amazon ECS

- Managed containers
- Highly scalable, high performance


AWS Fargate

Amazon EKS

Amazon Elastic Compute Cloud (EC2)

Amazon Elastic Compute Cloud

AWSSOME DAY
ONLINE CONFERENCE




- ✓ Application Server
- ✓ Web Server
- ✓ Database Server
- ✓ Game Server
- ✓ Mail Server
- ✓ Media Server
- ✓ Catalog Server
- ✓ File Server
- ✓ Computing Server
- ✓ Proxy Server

What is Amazon EC2?

AWSSOME DAY
ONLINE CONFERENCE

Amazon EC2 Instances

- Pay as you go
- Broad selection of HW/SW
- Global hosting
- *Much more* (aws.amazon.com/ec2)



Instance Types

Families	Description	Example Use Cases
t2, m3, m4, m5	General Purpose Balanced Performance	Websites, web applications, Dev, code repos, micro services, business apps
c3, c4, c5, cc2	Compute Optimized High CPU Performance	Front-end fleets, web-servers, batch processing, distributed analytics, science and engineering apps, ad serving, MMO gaming, video-encoding
g2, p2	GPU Optimized High-end GPU	Amazon AppStream 2.0, video encoding, machine learning, high perf databases, science
r3, r4, r5, x1, cr1	Memory Optimized Large RAM footprint	In-memory databases, data mining
d2, i2, i3, hi1, hs1	Storage Optimized High I/O, High density	NAS, data warehousing, NoSQL

Choosing the Right Amazon EC2 Instances

AWSSOME DAY
ONLINE CONFERENCE

- EC2 Instance types are optimized for different use cases, workloads & come in multiple sizes. This allows you to optimally scale resources to your workload requirements.
- AWS utilizes Intel® Xeon® processors for EC2 Instances providing customers with high performance and value.
- Consider the following when choosing your instances: core count, memory size, storage size & type, network performance, I/O requirements & CPU technologies.
- **Hurry Up & Go Idle** - A larger compute instance can save you time and money, therefore paying more per hour for a shorter amount of time can be less expensive.



EC2 Instances Powered by Intel Technologies





AWSSOME DAY
ONLINE CONFERENCE

EC2 Instance Type	Compute Optimized		General Purpose			Memory Optimized			Storage Optimized		
	C5n	C5	M5	T3	T2	X1	X1e	R4	H1	I3	D2
Intel Processor	Xeon Platinum 8175M	Xeon Platinum 8175M	Xeon Platinum 8175M	Xeon Platinum 8175M	Xeon Family	Xeon E7 8880 v3	Xeon E7 8880 v3	Xeon E5 2686 v4	Xeon E5 2686 v4	Xeon E5 2686 v4	Xeon E5 2676 v3
Intel Processor Technology	Skylake	Skylake	Skylake	Skylake	Yes	Haswell	Haswell	Broadwell	Broadwell	Broadwell	Haswell
Intel AVX	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Intel AVX2	Yes	Yes	Yes	Yes	-	Yes	Yes	Yes	Yes	Yes	Yes
Intel AVX-512	Yes	Yes	Yes	Yes	-	-	-	-	-	-	-
Intel Turbo Boost	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Storage	EBS-only	EBS-only	EBS-only	EBS-only	EBS-only	SSD EBS-Opt	SSD EBS-Opt	-	HDD	SSD	HDD

C5n: Compute Optimized Instances

AWSSOME DAY
ONLINE CONFERENCE


- Featuring Intel Xeon Scalable processors
- 100 Gbps network bandwidth on largest instance sizes
- 25 Gbps peak bandwidth on smaller instance sizes
- 33% Increased memory footprint over C5 instances



Faster analytics and
big data workloads



Lower costs for
network-bound workloads



All of the elasticity, security,
and scalability of AWS



R5: Memory-Optimized Instances

AWSSOME DAY
ONLINE CONFERENCE

- 2.5 GHz **Intel Xeon Scalable** processors (Skylake)
- **Memory-optimized** instances with **8:1 GiB to vCPU**
- Up to **25 Gbps NW bandwidth**
- R5d instances include up to **3.6 TB of local NVMe SSD**



r5.large
16 GiB
2 vCPU

6 sizes
● ● ●

r5.24xlarge
768 GiB
96 vCPU

In-memory caches



High performance databases



Big data analytics



R5a: Now available with
AMD EPYC 7000 processor



R5.metal Bare Metal instances
coming soon on **Intel Xeon Scalable** processors

z1d: High Frequency for Specialized Workloads

AWSSOME DAY
ONLINE CONFERENCE

- High Frequency instances with custom **Intel Xeon Scalable** processors running at sustained **4 GHz** all core turbo
- **8:1 GiB to vCPU ratio**
- Up to **25 Gbps network bandwidth** and up to **1.8 TB** of local NVMe storage


z1d.large	
16 GiB	
2 vCPU	

6 sizes



z1d.12xlarge	
384 GiB	
48 vCPU	

Electronic Design Automation



Relational databases



Gaming



z1d.metal Bare Metal instances coming soon


© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.



M5: Next-Gen General Purpose Instance

AWSSOME DAY
ONLINE CONFERENCE

- Powered by 2.5 GHz Intel Xeon Scalable Processors (Skylake)
- New larger instance size—m5.24xlarge with 96 vCPUs and 384 GiB of memory (4:1 Memory:vCPU ratio)
- Improved network and EBS performance on smaller instance sizes
- Support for Intel AVX-512 offering up to twice the performance for vector and floating point workloads



T3: Next-Gen General Purpose Instance


AWSSOME DAY
ONLINE CONFERENCE

- Balance of compute, memory, and network
- **Baseline level of CPU performance** with the ability to **burst CPU usage when needed** at any time for as long as required
- Lowest cost instance at \$0.0052 per hour and up to 30% better price performance over T2 using **Intel Xeon Scalable Processors**

t3.nano
0.5 GiB
2 vCPU
Base perf 5%

7 sizes
● ● ●

t3.2xlarge
32 GiB
8 vCPU
Base perf 40%



With T3 Unlimited bursting over baseline is only \$0.05 per vCPU-hour, averaged over 24 hours



Elastic Load Balancing (ELB)


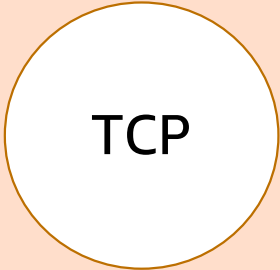
Introduction to Elastic Load Balancing

AWSSOME DAY
ONLINE CONFERENCE

Managed load balancing service
Distributes loads between instances

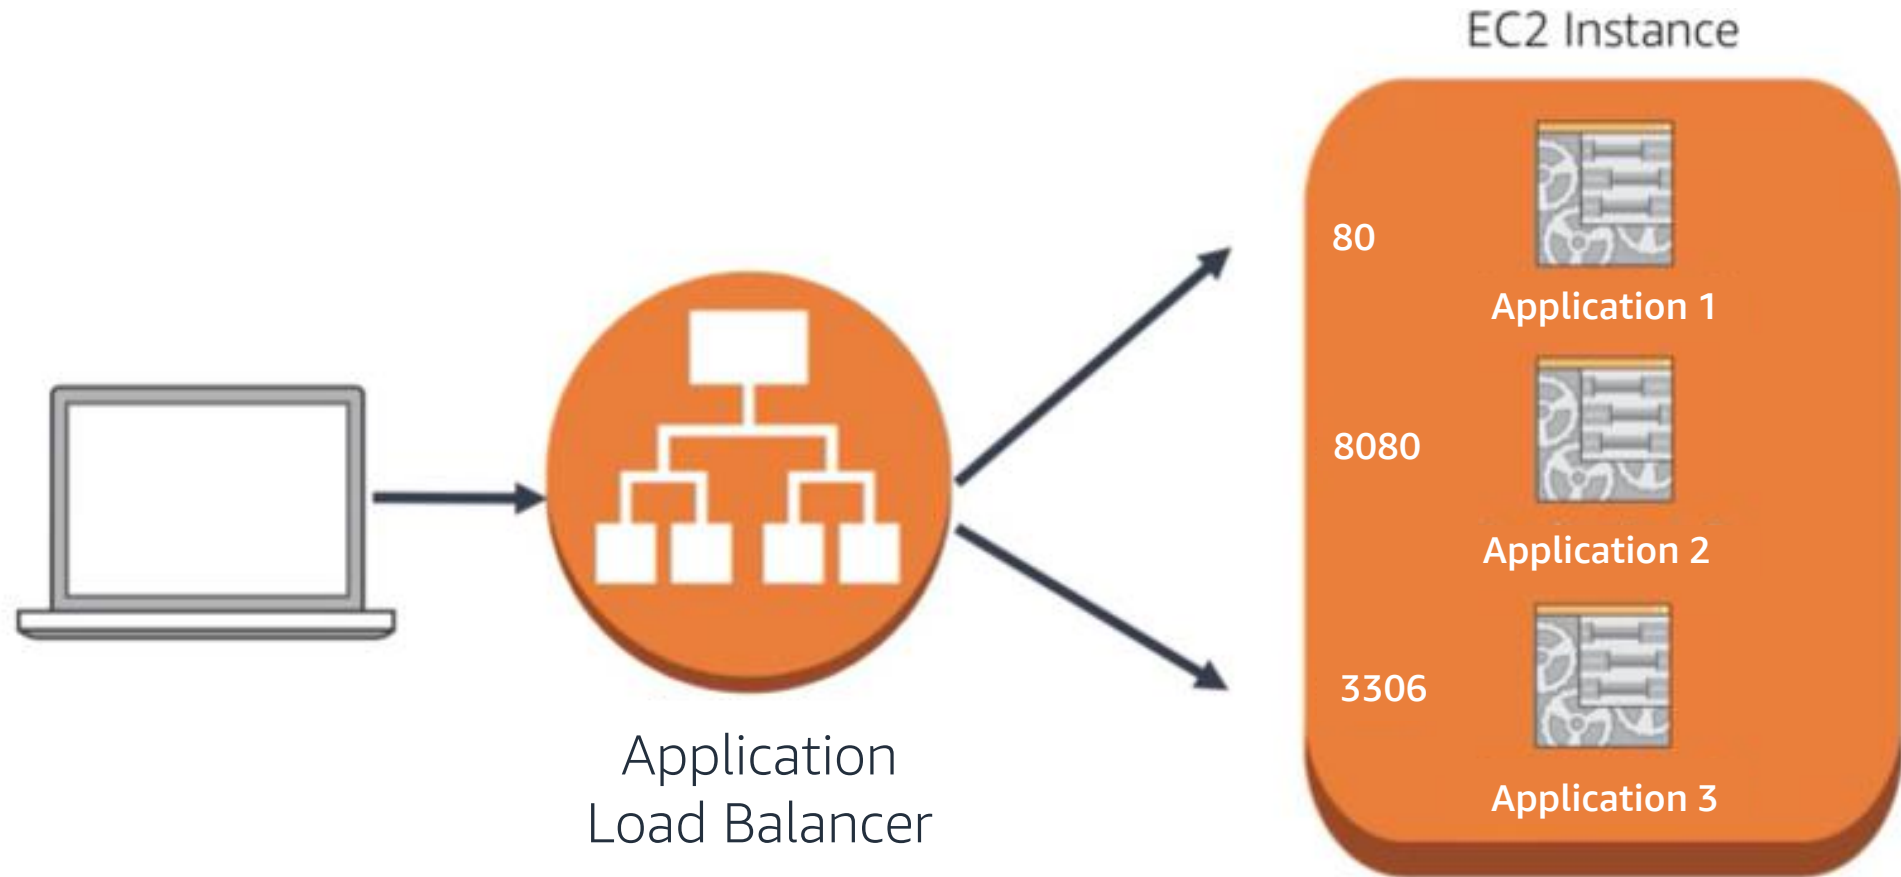
Elastic Load Balancing Products

AWSSOME DAY
ONLINE CONFERENCE

Application Load Balancer (ALB)	Network Load Balancer (NLB)	Classic Load Balancer (CLB)
		PREVIOUS GENERATION for HTTP, HTTPS, and TCP
<ul style="list-style-type: none">• Flexible application management• Advanced load balancing of HTTP and HTTPS traffic• Operates at the request level (Layer 7)	<ul style="list-style-type: none">• Extreme performance and static IP for your application• Load balancing of TCP traffic• Operates at the connection level (Layer 4)	<ul style="list-style-type: none">• Existing application that was built within the EC2-Classic network• Operates at both the request level and connection level

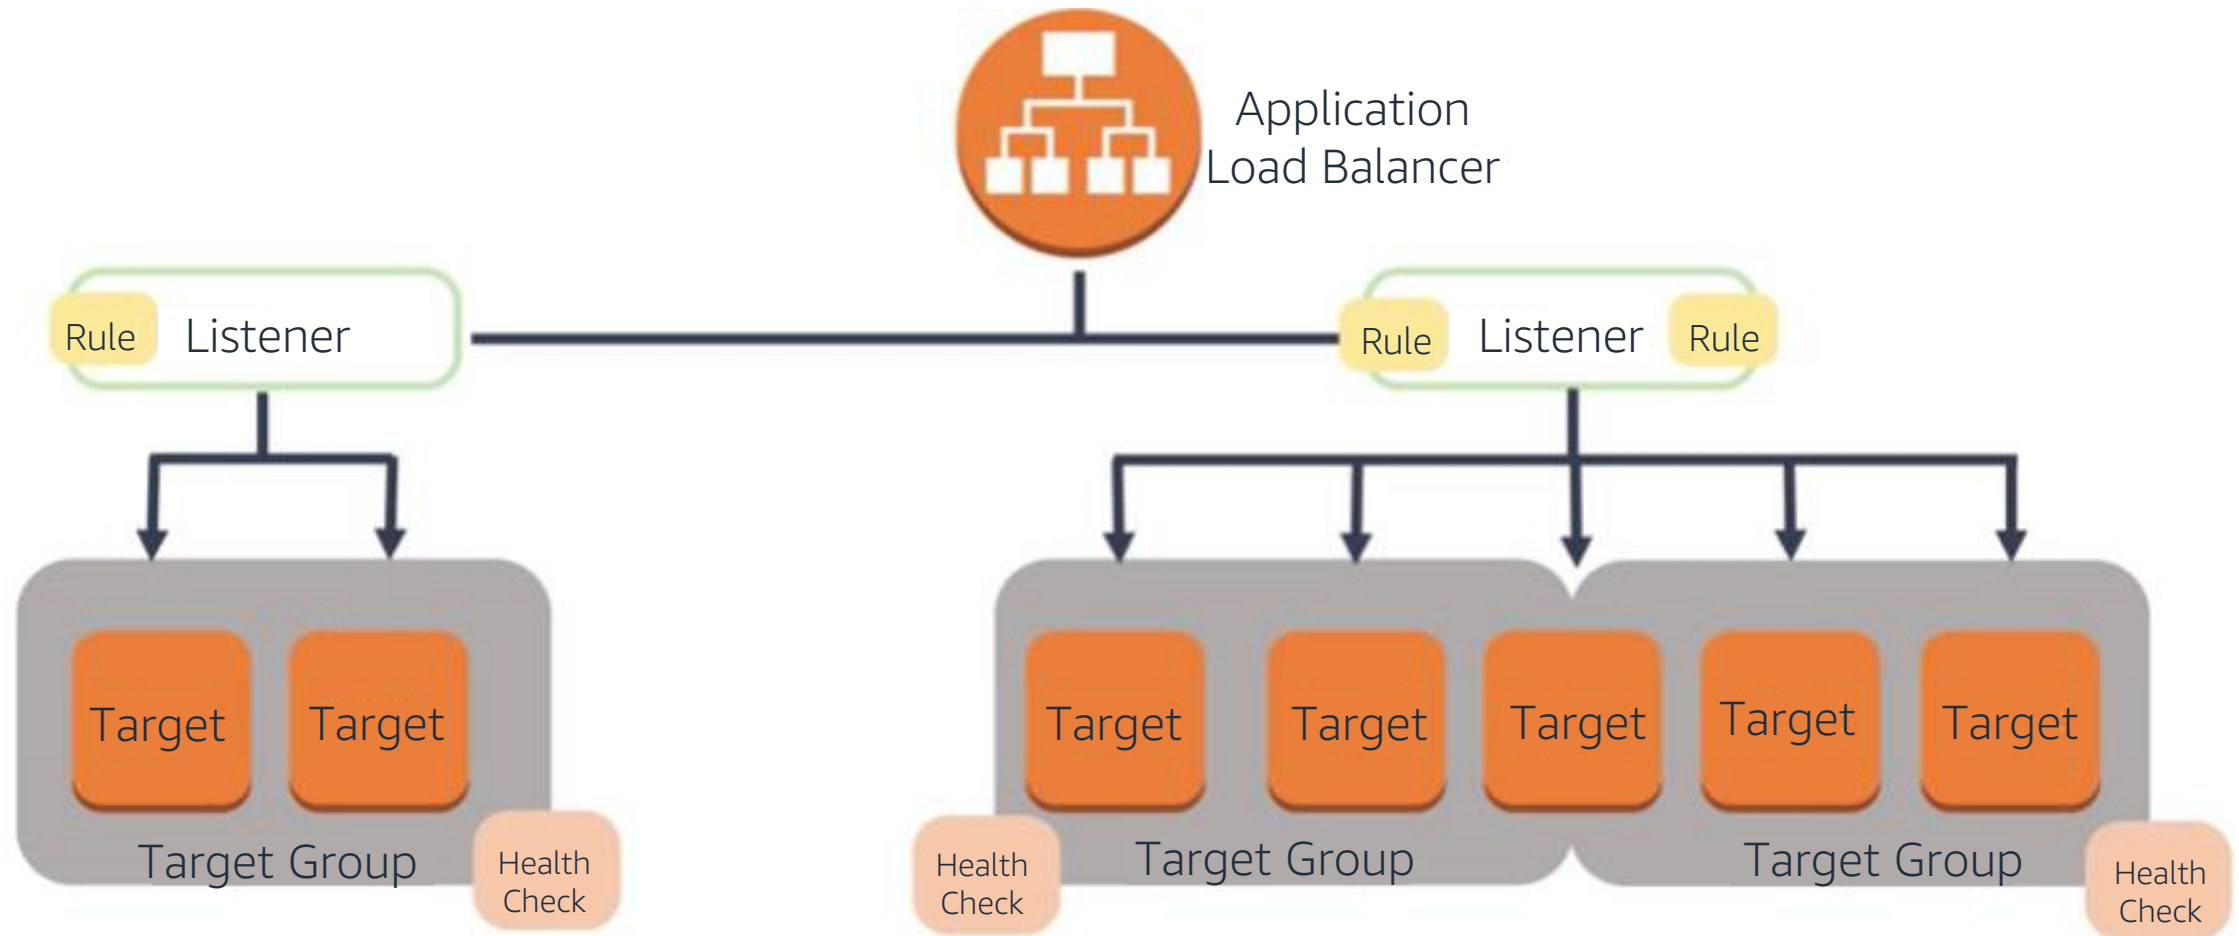
Application Load Balancer Use Cases

AWSSOME DAY
ONLINE CONFERENCE



Application Load Balancer Use Cases

AWSSOME DAY
ONLINE CONFERENCE



Network Load Balancer Use Cases

AWSSOME DAY
ONLINE CONFERENCE

Sudden and volatile traffic patterns

Single static IP address per Availability Zone


Ideal for applications that require extreme performance

AWS Auto Scaling

What Is AWS Auto Scaling?

AWSSOME DAY
ONLINE CONFERENCE

Helps you verify that you have the correct number of Amazon EC2 instances available to handle the load for your application





Monitoring Resource Performance

AWSSOME DAY
ONLINE CONFERENCE

Amazon CloudWatch to monitor performance
AWS Auto Scaling to add or remove Amazon EC2 instances

Capacity Management

AWSSOME DAY
ONLINE CONFERENCE



Critical Questions

AWSSOME DAY
ONLINE CONFERENCE

How can I make sure that my workload has enough Amazon EC2 resources to meet fluctuating performance requirements?

Scalability


How can Amazon EC2 resource provisioning to occur on-demand?

Automation


Scaling Out and Scaling In

AWSSOME DAY
ONLINE CONFERENCE


Elastic Load
Balancing



Base Configuration



Scaling Out
Launch Instances



Scaling In
Terminate Instances

Auto Scaling Components

AWSSOME DAY
ONLINE CONFERENCE

Launch Configuration
Auto Scaling groups
Auto Scaling Policy

Auto Scaling Components

AWSSOME DAY
ONLINE CONFERENCE

Launch Configuration: *What will be scaled?*

Launch settings

- AMI
- Instance type
- Security groups
- Roles

Auto Scaling Components

AWSSOME DAY
ONLINE CONFERENCE

Auto Scaling Group: *Where will it take place?*

Deployment settings

- VPC and subnets
- Load balancer
- Minimum instances
- Maximum instances
- Desired capacity

Auto Scaling Components

AWSSOME DAY
ONLINE CONFERENCE


Auto Scaling Policy: *When will it take place?*

Policy settings

- Scheduled
- On-demand
- Scale-out policy
- Scale-in policy

Dynamic Auto Scaling

AWSOME DAY
ONLINE CONFERENCE



© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Amazon CloudWatch Alarms for Auto Scaling

AWSSOME DAY
ONLINE CONFERENCE



Whenever: CPUUtilization

is:

for: consecutive period(s)

AutoScaling Action

Delete

Whenever this alarm:

From resource type:

From the:


Take this action:

Summary

Created

- A launch configuration
- Auto Scaling group
- Auto Scaling policy

Triggered Auto Scaling



Amazon Elastic Block Store (EBS)

Amazon EBS Volumes

AWSSOME DAY
ONLINE CONFERENCE

Characteristics

- Persistent and customizable block storage for Amazon EC2 instances
- HDD and SSD types
- Use Snapshots for backups
- Easy and transparent encryption
- Elastic

Amazon EBS Volumes

AWSSOME DAY
ONLINE CONFERENCE

Availability

- Durable and automatically replicated

Drive Types

- Storage that best fits your needs
- Magnetic or SSD
- Performance and price requirements

Snapshots


- Point-in-time snapshots
- Recreate a new volume at any time

Encryption

- Encrypted EBS volumes
- No additional cost

Elasticity

- Increase capacity
- Change to different types



Features

- Persistent and customizable block storage for Amazon EC2 instances
- HDD and SSD types
- Replicated in the same Availability Zone
- Easy and transparent encryption
- Elastic volumes
- Back up using snapshots

Amazon Simple Storage Service (S3)

Amazon S3

AWSSOME DAY
ONLINE CONFERENCE

- **Features**


- Fully managed cloud storage service
- Rich security controls

- **Functionality**

- Store virtually unlimited number of objects
- Access any time, from anywhere

Getting Started with Amazon S3

AWSOME DAY
ONLINE CONFERENCE



Example - <http://my-bucket-name.s3.amazonaws.com/welcome.mp4>

Access the Data Anywhere

AWSSOME DAY
ONLINE CONFERENCE

AWS Management console

AWS CLI

AWS SDKs

Common Use Cases

AWSSOME DAY
ONLINE CONFERENCE

Storing application assets
Static web hosting
Backup and disaster recovery (DR)
Staging area for big data

Summary

- Fully managed cloud storage service
- Store virtually unlimited number of objects
- Access any time, from anywhere
- Rich security controls
- Common use cases

Let's take a look at an Amazon S3 Demo

Amazon S3 Demo

Amazon Relational Database Service (RDS)

Challenges of Relational Databases

AWSSOME DAY
ONLINE CONFERENCE

Server maintenance and energy footprint

Software installation and patches

Database backups and high availability

Limits on scalability


Data security

OS install and patches

Amazon RDS

AWSSOME DAY
ONLINE CONFERENCE

Managed service that sets up and operates a relational database in the Cloud



© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Amazon RDS

AWSSOME DAY
ONLINE CONFERENCE

Customer manages

- Application Optimization
- Database schema
- Data

AWS manages

- OS installation and patches
- Database software installation and patches
- Database backups
- High availability
- Scaling
- Power, rack and stack
- Server maintenance

© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Amazon RDS DB Instances

AWSSOME DAY
ONLINE CONFERENCE

Amazon
RDS



RDS DB master
instance

DB Instance Class

- CPU
- Memory
- Network Performance

DB Instance Storage

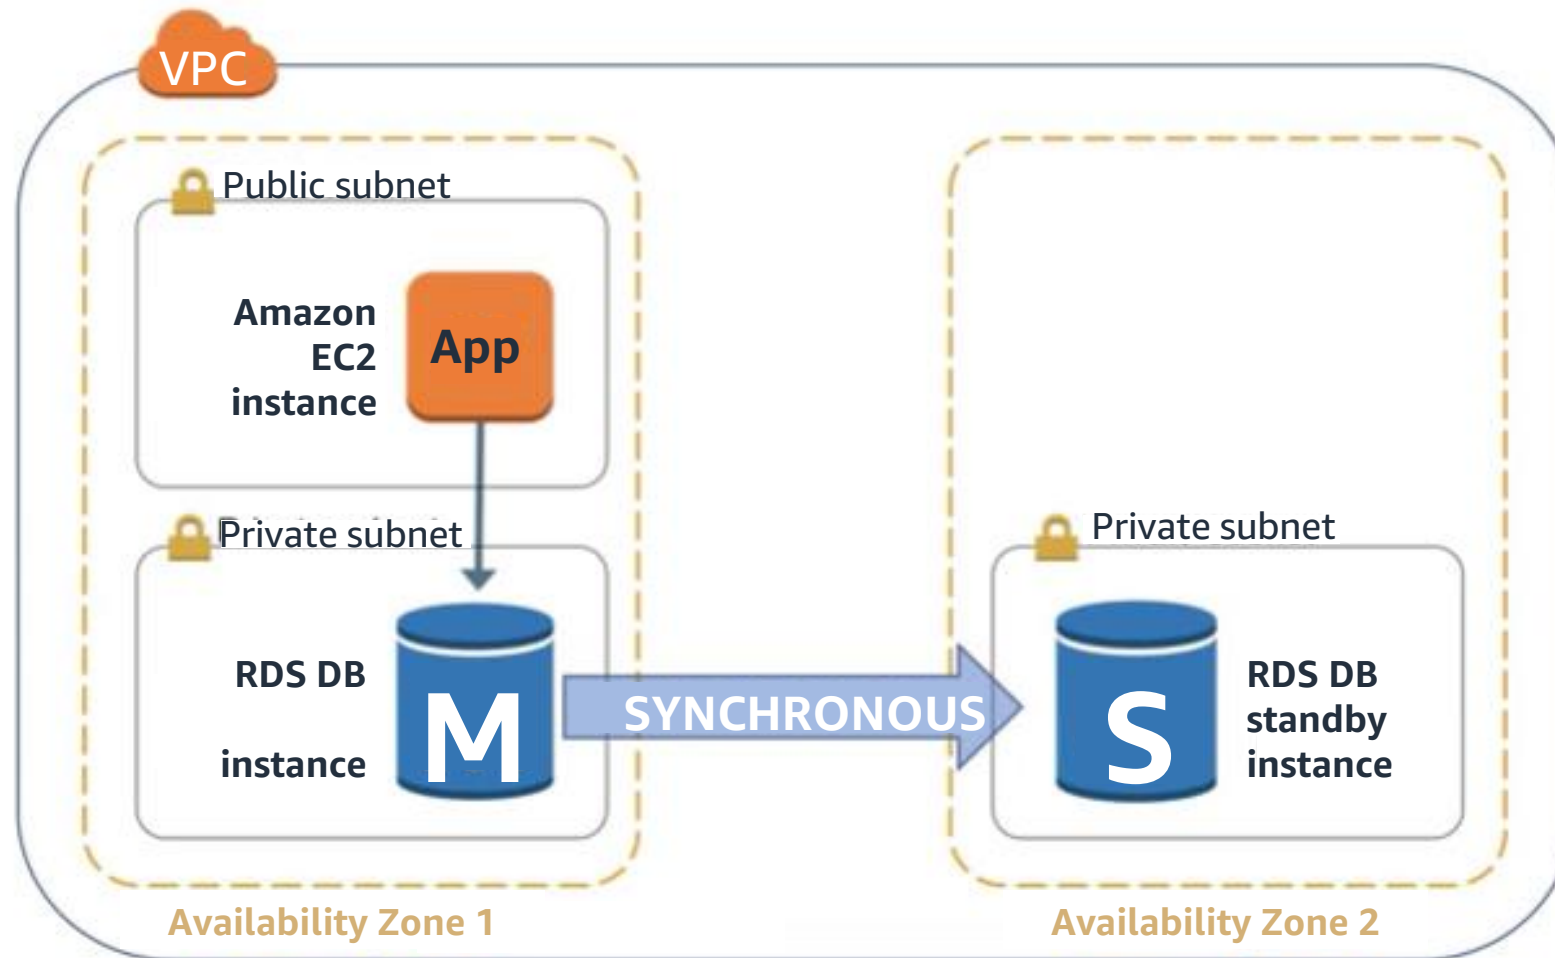
- Magnetic
- General Purpose (SSD)
- Provisioned IOPS



DB Engines

High Availability with Multi-AZ

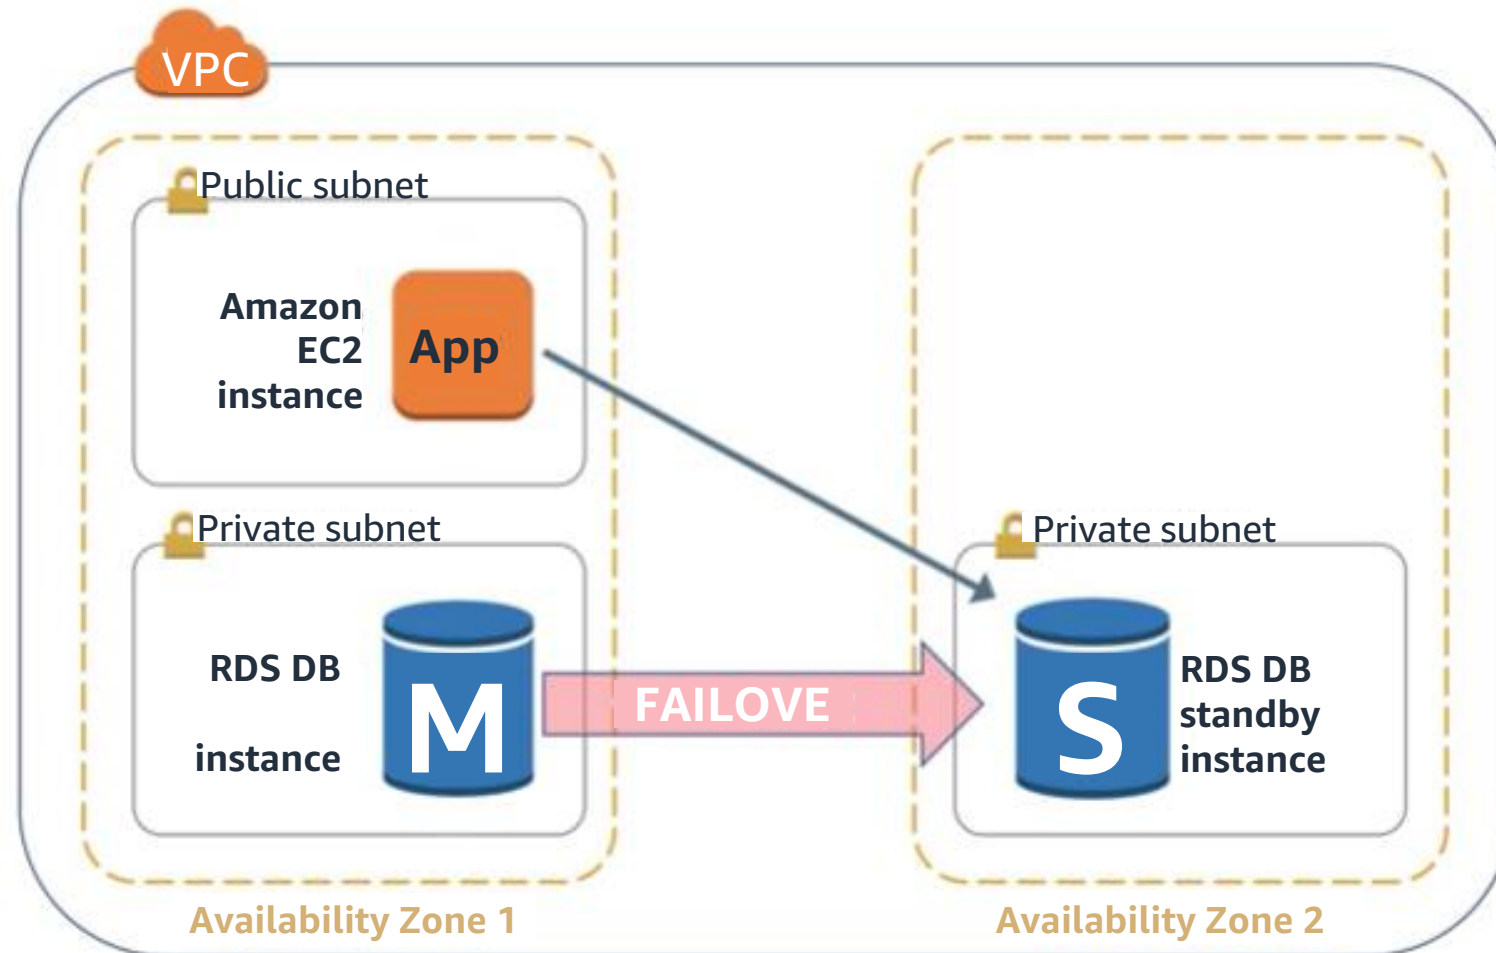
AWSOME DAY
ONLINE CONFERENCE



© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.

High Availability with Multi-AZ

AWSSOME DAY
ONLINE CONFERENCE



© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Amazon RDS Read Replicas


AWSSOME DAY
ONLINE CONFERENCE

Features

- Asynchronous replication
- Promote to master if necessary

Functionality

- Read-heavy database workloads
- Offload read queries



Summary

AWSSOME DAY
ONLINE CONFERENCE

Highly scalable

High performance

Easy to administer

Available and durable

Secure and compliant

Amazon DynamoDB

What Is Amazon DynamoDB?

AWSSOME DAY
ONLINE CONFERENCE

NoSQL database tables

Virtually unlimited storage

Items may have differing attributes

Low-latency queries

Scalable read/write throughput

Common Use Cases

AWSSOME DAY
ONLINE CONFERENCE

Web

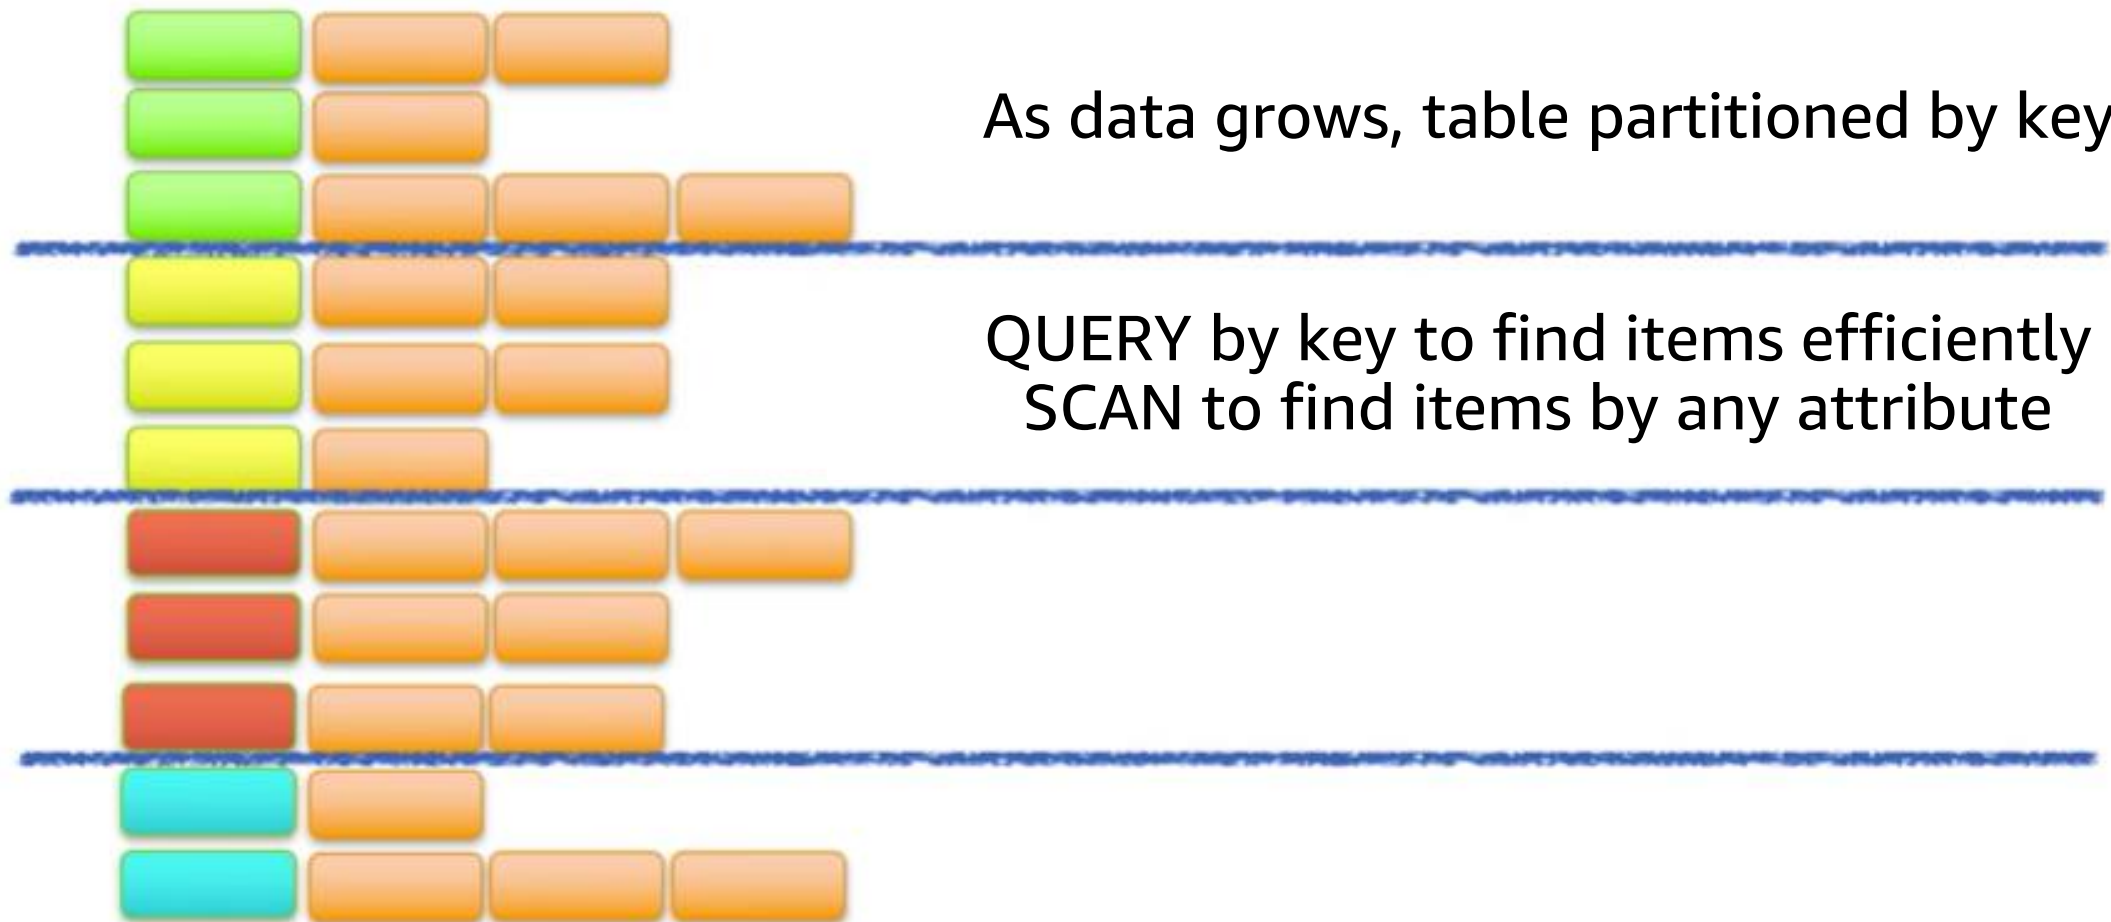
Mobile apps

Internet of Things

Ad tech

Gaming


Partitioning



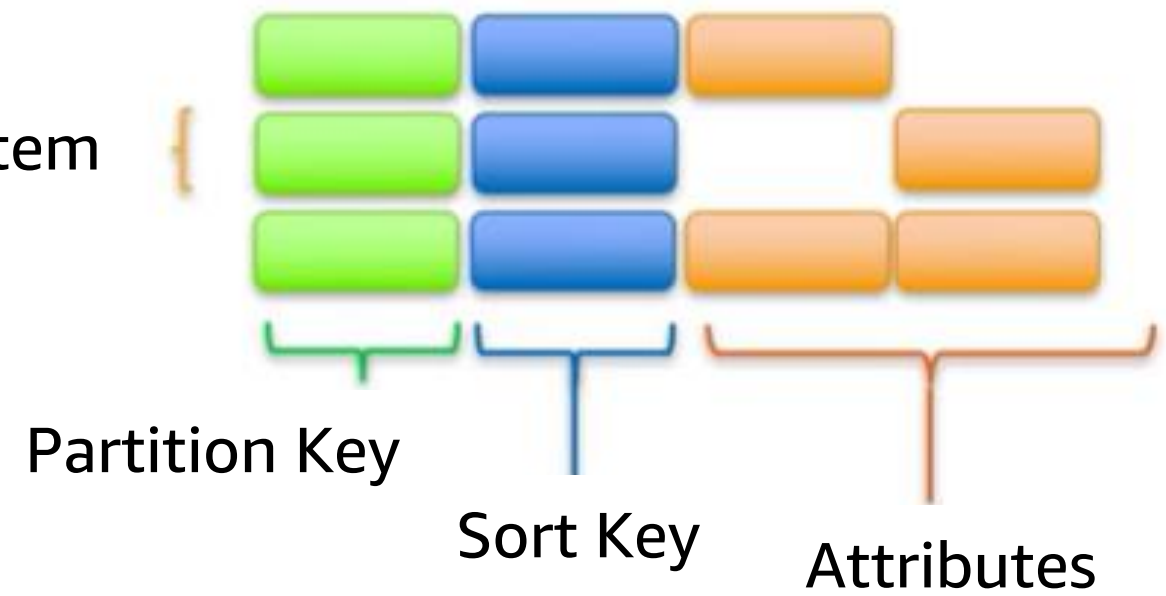
Items in a Table Must Have a Key

AWSSOME DAY
ONLINE CONFERENCE

Single Key



Compound Key



Managed NoSQL database service

Data store for applications

- Store large amounts of data
- Support high request volume
- Require low-latency query performance

End of Module 2

Test Your Knowledge